


Report on the Perceptron Algorithm

Model Architecture : The Perceptron is the simplest type of neural network. It consists of:

- **Input Layer:** Takes input features.
- **Linear Combination:** Computes a weighted sum of the inputs and adds a bias term.
- **Activation Function:** A step function determines the binary output.



Vector Representation of Data

- **Inputs:** $\mathbf{x} = [x\{1\}, x\{2\}, \dots, x\{n\}]$
- **Weights:** $\mathbf{w} = [w\{1\}, w\{2\}, \dots, w\{n\}]$
- **Bias:** b
- **Output:** $y \in \{0, 1\}$

Mathematical Formulation

$$z = \mathbf{w}^T \mathbf{x} + b = \sum_{i=1}^n w_i x_i + b$$

1. **Linear Combination:**

$$f(z) = \begin{cases} 1 & \text{if } z > 0 \\ 0 & \text{if } z \leq 0 \end{cases}$$

2. **Activation Function (Step Function):**

$$L(y, \hat{y}) = \frac{1}{2}(y - \hat{y})^2$$

3. **Loss Function:** To measure prediction error, we use:

Prediction Calculation

The Perceptron predicts output (\hat{y}) based on the following steps:

$$L(y, \hat{y}) = \frac{1}{2}(y - \hat{y})^2$$

1. Compute z , the weighted sum of inputs plus bias:

$$\begin{aligned} \frac{\partial L}{\partial w_i} &= -(y - \hat{y})x_i \\ \frac{\partial L}{\partial b} &= -(y - \hat{y}) \end{aligned}$$

2. Apply the activation function $f(z)$:

3. Update Rules:

- Update weights: $w_i = w_i - \eta \frac{\partial L}{\partial w_i}$

- Update bias: $b = b - \eta \frac{\partial L}{\partial b}$

Here, η is the learning rate.


Summary

- Perceptron** is a linear classifier that calculates a weighted sum of inputs, applies a step function, and updates its weights using gradient descent.
- It learns a decision boundary to separate two classes in the feature space.

Report on Logistic Regression

Model Architecture: Logistic Regression is a binary classification algorithm that applies the sigmoid function to a linear combination of inputs. The structure is as follows:

- Input Layer:** Takes input features.
- Linear Combination:** Computes a weighted sum of the inputs and adds a bias term.
- Activation Function (Sigmoid):** Converts the linear combination into a probability value.
- Output Layer:** Produces the predicted probability of the positive class.



Vector Representation of Data

- Inputs:** $\mathbf{x} = [x_1, x_2, \dots, x_n]$
- Weights:** $\mathbf{w} = [w_1, w_2, \dots, w_n]$
- Bias:** b
- Output (Probability):** $\hat{y} \in [0, 1]$

Mathematical Formulation

1. Linear Combination:

The linear combination of inputs is computed as:

$$z = \mathbf{w}^T \mathbf{x} + b = \sum_{i=1}^n w_i x_i + b$$

2. Activation Function (Sigmoid):

$$f(z) = \frac{1}{1 + e^{-z}}$$

The sigmoid function transforms z into a probability value:

3. **Loss Function (Binary Cross-Entropy):** To quantify the difference between predicted (\hat{y}) and true values (y), we use the binary cross-entropy loss:

$$L(y, \hat{y}) = -[y \log(\hat{y}) + (1 - y) \log(1 - \hat{y})]$$

Prediction Calculation

Logistic Regression predicts the output (\hat{y}) in the following steps:

1. Compute the linear combination:

$$z = \mathbf{w}^T \mathbf{x} + b$$

2. Apply the sigmoid function to z to get the probability:

$$\hat{y} = f(z) = \frac{1}{1 + e^{-z}}$$

Gradient Descent Algorithm

The objective of the gradient descent algorithm is to minimize the binary cross-entropy loss $L(y, \hat{y})$ by iteratively updating the weights and bias.

$$J(\mathbf{w}, b) = \frac{1}{m} \sum_{i=1}^m L(y^{(i)}, \hat{y}^{(i)})$$

1. **Objective:** Minimize:

where m is the number of training samples.

2. **Gradient Formulas:** The gradients of the loss with respect to the weights and bias are:

$$\frac{\partial J}{\partial w_i} = \frac{1}{m} \sum_{j=1}^m (\hat{y}^{(j)} - y^{(j)}) x_i^{(j)}$$

- For weight w_i :

$$\frac{\partial J}{\partial b} = \frac{1}{m} \sum_{j=1}^m (\hat{y}^{(j)} - y^{(j)})$$

- For bias b :

3. **Weight and Bias Updates:** Using the gradients, the weights and bias are updated as follows:

$$w_i = w_i - \eta \frac{\partial J}{\partial w_i}$$

- Weight update:

$$b = b - \eta \frac{\partial J}{\partial b}$$

- Bias update:

Here, η is the learning rate.

Summary

- Logistic Regression is a fundamental binary classifier that transforms a linear combination of inputs into a probability using the sigmoid function.
- The model minimizes the binary cross-entropy loss via gradient descent to iteratively update weights and bias.
- Its ability to output probabilities makes it widely applicable to binary classification tasks.


Report on Multilayer Perceptron (MLP)

Model Architecture: The Multilayer Perceptron (MLP) is a feedforward neural network consisting of one or more hidden layers between the input and output layers. Each neuron performs a linear combination of its inputs, applies an activation function, and passes the output to the next layer.

Architecture:

- **Input Layer:** Accepts the input features.
- **Hidden Layers:** Each layer consists of neurons with weights, biases, and nonlinear activation functions.
- **Output Layer:** Computes predictions (probabilities or class scores).

Diagram:



Vector Representation of Data

- **Inputs:** $x=[x_1, x_2, \dots, x_n]$
- **Weights (Layer L):** $W[l]$ (matrix of weights for layer L)
- **Bias (Layer L):** $b[l]$ (vector of biases for layer L)
- **Outputs (Activations for Layer L):** $a[l]$
- **Final Output:** \hat{y} (predicted value)

Mathematical Formulation

1. **Linear Combination (Layer L):** For each neuron in a layer:

$$z^{[l]} = \mathbf{W}^{[l]} \mathbf{a}^{[l-1]} + \mathbf{b}^{[l]}$$

where $\mathbf{a}^{[L-1]}$ is the activation output of the previous layer (or input data if $L=1$).

2. **Activation Function:** Apply a nonlinear function to $z^{[L]}$ to introduce nonlinearity.

Common activation functions:

- **ReLU:** $f(z) = \max(0, z)$

- **Sigmoid:** $f(z) = \frac{1}{1+e^{-z}}$

- **Tanh:** $f(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$

$$\mathbf{a}^{[l]} = f(z^{[l]})$$

Activations:

3. **Loss Function:** The loss function quantifies the error between predicted values \hat{y} and true values y . For binary classification:

$$L(y, \hat{y}) = -\frac{1}{m} \sum_{i=1}^m \left[y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}) \right]$$

$$L(y, \hat{y}) = \frac{1}{m} \sum_{i=1}^m \left(y^{(i)} - \hat{y}^{(i)} \right)^2$$

For regression, a common choice is the Mean Squared Error:

Prediction Calculation

1. **Forward Propagation:** Predictions are calculated by propagating inputs through all layers:

- Compute $z^{[l]}$: $z^{[l]} = \mathbf{W}^{[l]} \mathbf{a}^{[l-1]} + \mathbf{b}^{[l]}$

- Compute activations: $\mathbf{a}^{[l]} = f(z^{[l]})$

2. **Final Output:** At the output layer (L): $\hat{y} = \mathbf{a}^{[L]}$

Gradient Descent Algorithm

Gradient Descent minimizes the loss $L(y, \hat{y})$ by updating weights $\mathbf{W}^{[l]}$ and biases $\mathbf{b}^{[l]}$ using the gradients of the loss.

1. **Backward Propagation:** Gradients are computed for each layer starting from the output layer (using the chain rule).

○ Output layer gradient: $\delta^{[L]} = \frac{\partial L}{\partial z^{[L]}} = \hat{y} - y$

○ Hidden layer gradients (for $l < L$): $\delta^{[l]} = \left(\mathbf{W}^{[l+1]T} \delta^{[l+1]} \right) \odot f'(z^{[l]})$

where $f'(z^{[l]})$ is the derivative of the activation function.

2. Gradient Formulas:

○ Gradients of weights: $\frac{\partial L}{\partial \mathbf{W}^{[l]}} = \delta^{[l]} \mathbf{a}^{[l-1]T}$

○ Gradients of biases: $\frac{\partial L}{\partial \mathbf{b}^{[l]}} = \delta^{[l]}$

3. Update Rules:

○ Update weights: $\mathbf{W}^{[l]} = \mathbf{W}^{[l]} - \eta \frac{\partial L}{\partial \mathbf{W}^{[l]}}$

○ Update biases: $\mathbf{b}^{[l]} = \mathbf{b}^{[l]} - \eta \frac{\partial L}{\partial \mathbf{b}^{[l]}}$

Here, η is the learning rate.

Summary

- **Multilayer Perceptron (MLP):** A neural network with one or more hidden layers that can model complex, nonlinear relationships.
- **Forward Propagation:** Calculates the predicted values (\hat{y}) layer by layer.
- **Backward Propagation:** Updates weights and biases by propagating gradients backward through the network.
- **Optimization:** Uses gradient descent to minimize the loss function.